



Your agents are waiting on **you.**

One inbox for every coding agent you run.

PRE-SEED

USD 100K SAFE, USD 1M POST

VILNIUS



◆ NEEDS YOU 3 WORKING 5

ILLUSTRATION, RANKED BY WAIT

Oc	Upgrade the payment SDK to v5 OpenCode	● NEEDS YOU	34m 12s
Cc	Add rate limiting to /api/login Claude Code	● NEEDS YOU	8m 48s
Mu	Bump the Node image to 22 Muse	● NEEDS YOU	1m 40s
Cx	Fix the flaky checkout test Codex	● WORKING	
Cu	Write the README quickstart Cursor	● WORKING	
Pi	Port the settings page to the form kit Pi	● WORKING	
Gk	Split the orders service tests Grok	● WORKING	
Ds	Rename the billing env vars DeepSeek Harness	● WORKING	

ŽYGIMANTAS BAGDZEVIČIUS / LUKAS KVERAGA / ERNESTAS RIMKEVIČIUS

— THE PROBLEM

An agent that asks a question at 10:02 and is answered at 10:25 **did nothing for 23 minutes.**


ILLUSTRATION OF ONE WAIT. NO TOOL MEASURES THIS ACROSS VENDORS TODAY.



10:25 +23 min
IDLE, BLOCKED

HIGH-AI TEAMS VS THEIR OWN BASELINE

FAROS AI, 10,000+ DEVS, 1,255 TEAMS



The natural bottleneck on all of this is how fast I can review the results.

SIMON WILLISON, OCT 2025

More evidence (METR, Stack Overflow, DORA, Atlassian) is in the appendix.

— IT GETS WORSE WITH EVERY TOOL

70%

of 906 surveyed engineers use 2-4 AI coding tools. 15% use 5 or more (Pragmatic Engineer, 2026).

Four tools, four queues. Nobody sees the whole queue.




ILLUSTRATION: FOUR AGENT TOOLS, FOUR SEPARATE QUEUES

Claude Code's own docs now recommend running several sessions in parallel, each in its own git worktree.

Anthropic added weekly limits because some users ran Claude Code "24/7".

Cursor promised to "refund any unexpected charges" after its June 2025 pricing change.

AI coding revenue went vertical.

Parallel sessions are now in the vendors' own docs, and most engineers mix 2-4 tools. On a team, nobody sees all the agents, so the wait grows with every one.

4%

of public GitHub commits authored by Claude Code, an estimate (analysis cited by Anthropic, Feb 2026)

2M+

weekly Codex users, 5x since the start of 2026 (Reuters, Mar 2026)

CAPITAL IS MOVING IN:
TWO AGENT-WORKBENCH
ROUNDS IN FOUR MONTHS

MAR 2026

Conductor USD 22M

Series A (USD 24M raised in total), runs agents inside its app


JUL 2026

Superset USD 11.5M

seed led by USV, an agentic IDE

ANNUALISED RUN-RATE REVENUE, USD BILLIONS, ONE SCALE

LATEST FIGURE PER VENDOR, FEB-SEP 2026



THE PRODUCT

Every agent CLI you run in one ranked list. Oldest wait on top.

- 01

Deploy

One sheet starts an agent, each on its own branch.
- 02

Pinged


Claude Code, Codex, Cursor, plus 5 CLIs in beta and custom ones.
- 03

Act

The question in plain words, one reply box. Review is the diff beside the list.
- 04

Ship

Merge through GitHub or locally. Every merge is a signed record.



Real screens from the build, demo mode (scripted agents, no model runs) / docs/inbox.md "The loop" / src/shared/attention.ts, tests/inbox.test.ts

— ONE-MINUTE DEMO

Asked to merged, one pane.

- 01

Deploy

Five branches, five agents
- 02

Pinged

Codex jumps to the top
- 03


Act

Answered in one box
- 04

Ship

Merged, wait signed

HOSTED DEMO ON REQUEST
27 S GIF FOR EMAIL



Recording: about a minute, silent, the real build in demo mode with scripted stand-in agents / tests/demo.test.ts, tests/demo-e2e.test.ts

Self-hosted, vendor-neutral, on your own machine.

Kipdeck itself is self-hosted: no Kipdeck cloud, no account. Your agents still talk to their own model APIs.



- INSTALLS FROM A LOCAL TARBALL TODAY; NPM AT M1
- NO PASSWORD ON YOUR OWN MACHINE
- NO KIPDECK CLOUD; AGENTS USE THEIR OWN APIS

PROOF OF EXECUTION

A working product, measured. No users yet.

FRESH START TO FIRST MERGED CHANGE

10.7s



4 CLICKS • MEASURED 2026-10-07, HEADLESS, STAND-IN AGENT, NO MODEL

3.4s

to the first agent at work,
1 click

1,281

test cases on the product
branch

19

security audit findings
closed before launch

Next, before revenue: publish, launch, and sign 5 design-partner teams that run 5+ agents a day.

— MARKET, BOTTOM-UP

11.3M developers use agents.

Professional developers
SlashData, early 2025

36.5M

x 31% actively use AI agents
Stack Overflow 2025 share; our estimate

11.3M

x USD 360 a seat a year = TAM
our team-tier price; at Conductor's USD 720: 8.1B

USD 4.1B

Daily agent users, 14.1% = SAM
Stack Overflow 2025; 5.1M x USD 360


USD 1.9B

Year-5 path, our assumption
100k free installs x 2% become teams x 10 seats
x USD 360


USD 7.2M ARR

The budget already exists: USD 5.5B+ a year already goes to AI coding tools (slide 4).

EACH DOT = 10,000 DEVELOPERS



Conductor and Superset help one developer run more agents. We stop your team keeping them waiting.



Distribution first. Then wait data no single vendor has.

- 01

Open-source installs

Single-player is MIT-licensed and free; public release at M1.
- 02

Cross-vendor wait data

Every merge already records agent, model, reviewer and wait. As teams adopt, that becomes data no single vendor can collect.
- 03

Team upgrade

The second person who needs the queue turns a free install into a paid team.

SIGNED MERGE RECORDS

ILLUSTRATION

#0412 ✓ signed

Codex

waited 4m

#0413 ✓ signed

Claude Code

waited 1m

#0414 ✓ signed

Cursor

waited 9m

MEDIAN WAIT ON A PERSON, BY AGENT, WEEK OVER WEEK

MOCK DATA

	W1	W2	W3	W4	W5
Claude Code	14m	9m	6m	3m	1m
Codex	22m	12m	8m	4m	2m
Cursor	11m	7m	4m	2m	1m
OpenCode	18m	13m	9m	5m	3m
Pi	9m	6m	3m	1m	1m

WAIT ☐ under 2m ☒ 2-10m ☐ over 10m

🔒 Self-hosted: the record sits on the team's machine, not in a Kipdeck cloud. The code is not the moat; installs and the wait data are.

Free for one engineer. Teams pay.


SINGLE PLAYER FreeMIT Inbox, local merge, signed log, wait metrics The distribution: one npx command, public at M1	TEAM TIERPRICE IS OUR ASSUMPTIONTEAM USD 30/seat/mo Shared queue, routing to the owner, team wait-time reports SSO, audit export. Not built yet: milestone M3
--	---

GO-TO-MARKET

01 Launch npm publish and Hacker News with the 60 s demo (M1).	02 5 design partners Founder-led, teams running 5+ agents a day; European and regulated first.	03 Bottom-up One engineer installs free; the team upgrades when a second person needs the queue.	04 Team tier Paid seats for design partners first, then self-serve.
---	---	---	--

No revenue yet. Paid design-partner pilots are possible (offer ready, none sold); that is design-partner revenue, not the business.

Engineers who have built and scaled products before.



Žygimantas Bagdzevičius
CEO / CTO • [LINKEDIN](#)

400k MAU, EXITED

Ex-Amazon engineer. Founded and exited a Lithuanian business that scaled to 400k monthly active users. Built Motored's whole product before outside money.



Lukas Kveraga
FOUNDING ENGINEER

Ex-Vinted

Lithuania's marketplace unicorn. Hands-on experience scaling a high-traffic two-sided marketplace.



Ernestas Rimkevičius
SOFTWARE ENGINEER • [LINKEDIN](#)


Pipelines

Backend, importer, listings pipeline and ops tooling (at Motored).

DISCLOSURE AND COMMITMENT

Today Žygimantas is also co-founder and CEO/CTO of Motored (Antler-backed, raising its own pre-seed), and Lukas and Ernestas are engineers there. Who goes full-time on Kipdeck and from when, the equity split and IP assignment to the Kipdeck company are agreed in writing before this round closes and shared in diligence.


Company: to be incorporated before close.




THE ASK

USD 100k

Post-money SAFE, USD 1M valuation cap: 10%.
Company to be incorporated before close. About USD 11k a month: 9 months of runway.



9 MONTHS TO 5 DESIGN PARTNERS AND THE FIRST PAYING TEAMS



NEEDS YOU 0			
0c	Upgrade the payment SDK to v5 OpenCode	✓ MERGED	waited 34m
cc	Add rate limiting to /api/login Claude Code	✓ MERGED	waited 8m
Mu	Bump the Node image to 22 Muse	✓ MERGED	waited 1m

Nobody waiting.
What this round proves: teams pay for one queue across vendors.

Žygimantas Bagdzevičius, Lukas Kveraga, Ernestas Rimkevičius

Tool by tool, and more evidence on the problem.

WHAT YOU NEED	VENDOR DASHBOARD	TASK BOARD	GITHUB PRS	TERMINALS, TMUX	KIPDECK
Every agent CLI in one list	Its own	If you file them	Once a PR exists	One pane at a time	Ranked by wait
Knows an agent is blocked on you	Its own	No	No	If you look	Top of the list
Answer in one box	In its app	No	No	Find the terminal	Every CLI
Review and merge before a PR	No	No	No	By hand	Or open one
Measures wait on a person	No	No	No	No	Per agent, model
Self-hosted	Varies (local CLI or vendor cloud)	Depends	Their cloud	Yes	Yes, MIT (public at M1)

METR: experienced devs were 19% slower with AI while believing they were 20% faster.

Stack Overflow 2025: 66% frustrated by "almost right" output; 3.1% highly trust AI.

DORA 2025: AI adoption has a negative relationship with delivery stability.

Atlassian DevEx, 3,500 devs: 50% lose 10+ hours a week to friction; 63% say leaders do not get it.