



Your agents are waiting on **you.**

One inbox for every coding agent you run.

PRE-SEED

USD 100K SAFE, USD 1M POST

VILNIUS

◆ NEEDS YOU 8		ILLUSTRATION, RANKED BY WAIT	
Oc	Upgrade the payment SDK to v5 OpenCode	● NEEDS YOU	10:00
Cc	Add rate limiting to /api/login Claude Code	● NEEDS YOU	08:48
Mu	Bump the Node image to 22 Muse	● NEEDS YOU	07:36
Cx	Fix the flaky checkout test Codex	● NEEDS YOU	06:24
Ds	Rename the billing env vars DeepSeek Harness	● NEEDS YOU	05:12
Cu	Write the README quickstart Cursor	● NEEDS YOU	04:00
Gk	Split the orders service tests Grok	● NEEDS YOU	02:48
Pi	Port the settings page to the form kit Pi	● NEEDS YOU	01:36

ŽYGIMANTAS BAGDZEVIČIUS / LUKAS KVERAGA / ERNESTAS RIMKEVIČIUS

THE PROBLEM

An agent that asks a question at 10:02 and is answered at 10:25 **did nothing for 23 minutes.**


ILLUSTRATION OF ONE WAIT. NO TOOL MEASURES THIS ACROSS VENDORS TODAY.



10:25 +23 min
IDLE, BLOCKED

HIGH-AI TEAMS VS THEIR OWN BASELINE

FAROS AI, 10,000+ DEVS, 1,255 TEAMS



The natural bottleneck on all of this is how fast I can review the results.

SIMON WILLISON, OCT 2025

More evidence (METR, Stack Overflow, DORA, Atlassian) is in the appendix.

— IT GETS WORSE WITH EVERY TOOL

70%

of 906 surveyed engineers use 2-4 AI coding tools. 15% use 5 or more (Pragmatic Engineer, 2026).

Four tools, four queues. Nobody sees the whole queue.




ILLUSTRATION: FOUR AGENT TOOLS, FOUR SEPARATE QUEUES

Claude Code's own docs now recommend running several sessions in parallel, each in its own git worktree.

Anthropic added weekly limits because some users ran Claude Code "24/7".

Cursor promised to "refund any unexpected charges" after its June 2025 pricing change.

AI coding revenue went vertical.

Parallel sessions are now in the vendors' own docs, and most engineers mix 2-4 tools. On a team, nobody sees all the agents, so the wait grows with every one.

4%

of public GitHub commits authored by Claude Code, an estimate (analysis cited by Anthropic, Feb 2026)

2M+

weekly Codex users, 5x since the start of 2026 (Reuters, Mar 2026)

CAPITAL IS MOVING IN:
THREE ROUNDS IN SEVEN MONTHS

FEB 2026

Entire USD 60M
seed, agent history in git

MAR 2026


Conductor USD 22M
Series A, parallel agents per dev

SEP 2026

Herdr USD 6M
seed, open-source agent runtime

ANNUALISED RUN-RATE REVENUE, USD BILLIONS, ONE SCALE

LATEST FIGURE PER VENDOR, FEB-SEP 2026



— THE PRODUCT


Every agent CLI you run in one ranked list. Oldest wait on top.

01 **Deploy**
One sheet starts an agent, each on its own branch.

02 **Pinged**
Claude Code, Codex, Cursor, plus 5 CLIs in beta and custom ones.

03 **Act**
The question in plain words, one reply box. Review is the diff beside the list.

04 **Ship**
Merge through GitHub or locally. Every merge is a signed record.





60-SECOND DEMO

Asked to merged, one pane.

The real build in demo mode: scripted stand-in agents, so it runs the same every time. A recording with live Claude Code and Codex sessions is next.

HOSTED DEMO ON REQUEST


30 S GIF FOR EMAIL



Engineers now run several coding agents at once. The bottleneck isn't the agents: it's how long they sit waiting on you.

Self-hosted, vendor-neutral, one command.

Kipdeck itself is self-hosted: no Kipdeck cloud, no account. Your agents still talk to their own model APIs.



- NPX KIPDECK: ONE COMMAND (LOCAL TARBALL TODAY; NPM AT M1)
- NO PASSWORD ON YOUR OWN MACHINE
- NO KIPDECK CLOUD; AGENTS USE THEIR OWN APIS

PROOF OF EXECUTION

A working product, measured. No users yet.

Next, before revenue: publish, launch, and sign 5 design-partner teams that run 5+ agents a day.

FRESH START TO FIRST MERGED CHANGE

10.7s

3.4s

to the first agent at work, 1 click

1 command


to install with npx; from a local tarball today, npm publish is M1

1,281


test cases on the product branch

19

security audit findings fixed in our fork



4 CLICKS • MEASURED 2026-10-07, HEADLESS, STAND-IN AGENT, NO MODEL



BEFORE: THE 3D PROOF OF CONCEPT → AFTER: THE INBOX, 8 DAYS LATER

measure.json (2026-10-07) / Test count: it/test calls in tests/, product/fundable c8119c0, 2026-10-07 / Audit: launch/report/comparison.md (2026-10-06)

— MARKET, BOTTOM-UP

11.3M developers use agents.

Professional developers

SlashData, early 2025

36.5M

x 31% actively use AI agents

Stack Overflow 2025 share; our estimate

11.3M

x USD 360 a seat a year = TAM

our team-tier price; at Conductor's USD 720: 8.1B

USD 4.1B

Daily agent users, 14.1% = SAM

Stack Overflow 2025; 5.1M x USD 360

USD 1.9B

Year-5 path, our assumption


100k free installs x 2% become teams x 10 seats

x USD 360


USD 7.2M ARR

The budget already exists: USD 5.5B+ a year already goes to AI coding tools (slide 4).

EACH DOT = 10,000 DEVELOPERS



Others see one vendor or one laptop. We are building the team's queue.



VENDOR DASHBOARDS

Claude Code's Agent view and Cursor's cloud agents each see their own agents. 70% of 906 surveyed engineers use 2-4 AI coding tools.

CONDUCTOR (USD 22M), HERDR (USD 6M)

Conductor sells a Teams plan but runs locally on a Mac, per developer. Herdr is an open-source runtime that shows each agent as working, blocked or idle, on one machine.

WHY WON'T HERDR ADD A TEAM TIER NEXT QUARTER?

It might. Our bet is the one number nobody reports: human wait time per agent and model, across vendors. We already log it on every merge. The team tier (shared queue, routing) is not built yet; it is milestone M3.

Distribution first. Then wait data no single vendor has.

- 01

Open-source installs

Single-player is MIT-licensed and free; public release at M1.
- 02


Cross-vendor wait data

Every merge already records agent, model, reviewer and wait. As teams adopt, that becomes data no single vendor can collect.
- 03

Team upgrade

The second person who needs the queue turns a free install into a paid team.

SIGNED MERGE RECORDS (ILLUSTRATION)



 Self-hosted: the record sits on the team's machine, not in a Kipdeck cloud.

Built on webdevcody/agent-office (MIT), so anyone can fork us too. What we added in 8 days: the ranking, inbox, question card, signed log, wait metrics, local merge, onboarding, attach and demo. The code is not the moat; installs and the wait data are.

Free for one engineer. Teams pay.


SINGLE PLAYER Free MIT Inbox, local merge, signed log, wait metrics The distribution: one npx command, public at M1	TEAM TIER (PRICE IS OUR ASSUMPTION) TEAM USD 30 /seat/mo Shared queue, routing to the owner, team wait-time reports SSO, audit export. Not built yet: milestone M3
--	--

GO-TO-MARKET

01 Launch npm publish and Hacker News with the 60 s demo (M1).	02 5 design partners Founder-led, teams running 5+ agents a day; European and regulated first.	03 Bottom-up One engineer installs free; the team upgrades when a second person needs the queue.	04 Team tier Paid seats for design partners first, then self-serve.
---	---	---	--

No revenue yet. Paid design-partner pilots are possible (offer ready, none sold); that is design-partner revenue, not the business.

Engineers who have built and scaled products before.



Žygimantas Bagdzevičius
CEO / CTO • [LINKEDIN](#)

400k MAU, EXITED

Ex-Amazon engineer. Founded and exited a Lithuanian business that scaled to 400k monthly active users. Built Motored's whole product before outside money.



Lukas Kveraga
FOUNDING ENGINEER

Ex-Vinted

Lithuania's marketplace unicorn. Hands-on experience scaling a high-traffic two-sided marketplace.



Ernestas Rimkevičius
SOFTWARE ENGINEER • [LINKEDIN](#)

Pipelines


Backend, importer, listings pipeline and ops tooling (at Motored).

DISCLOSURE AND COMMITMENT

Today Žygimantas is also co-founder and CEO/CTO of Motored (Antler-backed, raising its own pre-seed), and Lukas and Ernestas are engineers there. Who goes full-time on Kipdeck and from when, the equity split and IP assignment to the Kipdeck company are agreed in writing before this round closes and shared in diligence.

Company: to be incorporated before close.


PRODUCT-BRANCH COMMITS BY THE CEO/CTO WITH CODING AGENTS, 30 SEP - 7 OCT 2026




THE ASK

USD 100k

Post-money SAFE, USD 1M valuation cap: 10%.
Company to be incorporated before close. About USD 11k a month: 9 months of runway.



9 MONTHS TO 5 DESIGN PARTNERS AND THE FIRST PAYING TEAMS



✓ NEEDS YOU 0		NOBODY WAITING.	
Cx	Fix the flaky checkout test Codex	MERGED	00:00
CC	Add rate limiting to /api/login Claude Code	MERGED	00:00
Cu	Write the README quickstart Cursor	MERGED	00:00

What this round proves: teams pay for one queue across vendors. That is the case for a seed round. The long-term aim is the first fifty teams.

Žygimantas Bagdzevičius, Lukas Kveraga, Ernestas Rimkevičius

Tool by tool, and more evidence on the problem.

WHAT YOU NEED	VENDOR DASHBOARD	TASK BOARD	GITHUB PRS	TERMINALS, TMUX	KIPDECK
Every agent CLI in one list	Its own	If you file them	Once a PR exists	One pane at a time	Ranked by wait
Knows an agent is blocked on you	Its own	No	No	If you look	Top of the list
Answer in one box	In its app	No	No	Find the terminal	Every CLI
Review and merge before a PR	No	No	No	By hand	Or open one
Measures wait on a person	No	No	No	No	Per agent, model
Self-hosted	Varies (local CLI or vendor cloud)	Depends	Their cloud	Yes	Yes, MIT (public at M1)

METR: experienced devs were 19% slower with AI while believing they were 20% faster.

Stack Overflow 2025: 66% frustrated by "almost right" output; 3.1% highly trust AI.

DORA 2025: AI adoption has a negative relationship with delivery stability.

Atlassian DevEx, 3,500 devs: 50% lose 10+ hours a week to friction; 63% say leaders do not get it.