



Your agents are waiting on **you.**

One inbox for every coding agent you run.

PRE-SEED

USD 100K SAFE, USD 1M POST

VILNIUS



◆ NEEDS YOU 3 WORKING 5

ILLUSTRATION, RANKED BY WAIT


Oc	Upgrade the payment SDK to v5 OpenCode	● NEEDS YOU	34m 12s
Cc	Add rate limiting to /api/login Claude Code	● NEEDS YOU	8m 48s
Mu	Bump the Node image to 22 Muse	● NEEDS YOU	1m 40s
Cx	Fix the flaky checkout test Codex	● WORKING	
Cu	Write the README quickstart Cursor	● WORKING	
Pi	Port the settings page to the form kit Pi	● WORKING	
Gk	Split the orders service tests Grok	● WORKING	
Ds	Rename the billing env vars DeepSeek Harness	● WORKING	

ŽYGIMANTAS BAGDZEVIČIUS / LUKAS KVERAGA / ERNESTAS RIMKEVIČIUS

— THE PROBLEM

An agent that asks a question at 10:02 and is answered at 10:25 **did nothing for 23 minutes.**


ILLUSTRATION OF ONE WAIT. NO TOOL MEASURES THIS ACROSS VENDORS TODAY.



10:25 +23 min
IDLE, BLOCKED

HIGH-AI TEAMS VS THEIR OWN BASELINE

FAROS AI, 10,000+ DEVS, 1,255 TEAMS



The natural bottleneck on all of this is how fast I can review the results.

SIMON WILLISON, OCT 2025

More evidence (METR, Stack Overflow, DORA, Atlassian) is in the appendix.

— IT GETS WORSE WITH EVERY TOOL

70%

of 906 surveyed engineers use 2-4 AI coding tools. 15% use 5 or more (Pragmatic Engineer, 2026).

Four tools, four queues. Nobody sees the whole queue.




ILLUSTRATION: FOUR AGENT TOOLS, FOUR SEPARATE QUEUES

Claude Code's own docs now recommend running several sessions in parallel, each in its own git worktree.

Anthropic added weekly limits because some users ran Claude Code "24/7".

Cursor promised to "refund any unexpected charges" after its June 2025 pricing change.

AI coding revenue went vertical.

Parallel sessions are now in the vendors' own docs, and most engineers mix 2-4 tools. On a team, nobody sees all the agents, so the wait grows with every one.

4%

of public GitHub commits authored by Claude Code, an estimate (analysis cited by Anthropic, Feb 2026)

2M+

weekly Codex users, 5x since the start of 2026 (Reuters, Mar 2026)

CAPITAL IS MOVING IN:
TWO AGENT-WORKBENCH
ROUNDS IN FOUR MONTHS

MAR 2026

Conductor USD 22M


Series A (USD 24M raised in total), runs agents inside its app

JUL 2026


Superset USD 11.5M

seed led by USV, an agentic IDE

ANNUALISED RUN-RATE REVENUE, USD BILLIONS, ONE SCALE



LATEST FIGURE PER VENDOR, FEB-SEP 2026



— THE PRODUCT

Every agent CLI you run in one ranked list. Oldest wait on top.

01 Deploy

One sheet starts an agent, each on its own branch.

02 Pinged



Claude Code, Codex, Cursor, plus 5 CLIs in beta and custom ones.

03 Act

The question in plain words, one reply box. Review is the diff beside the list.

04 Ship

Merge through GitHub or locally. Every merge is a signed record.



signed: agent / model / reviewer / wait

60-SECOND DEMO

Asked to merged, one pane.

- 01

Deploy

Five branches, five agents
- 02

Pinged

Codex jumps to the top
- 03

Act

Answered in one box
- 04

Ship

Merged, wait signed

HOSTED DEMO ON REQUEST
30 S GIF FOR EMAIL

npx kipdeck --demo

DEMO DATA: SCRIPTED AGENTS, NO MODEL RUNS

Kipdeck

DEMO Run it for real from your clone: node bin/agent-office.js

Search agents

Ctrl K

1 waiting on you

median wait

0 merged today

+ Deploy agent

DL

Needs you 1

Cx

Fix the flaky checkout test

Needs an answer

waiting <1m

Answer

To review

Nothing to review.

Working 4

CC

Add rate limiting to /api/login

Edit: api/login.js

<1m

Cu

Write the README quickstart

Read: package.json

<1m

CC

Upgrade the payment SDK to v5

Bash: npm outdated acme-pay-sdk

<1m

Cx

Port the settings page to the form kit

Search: form kit usages

<1m

Cx

Fix the flaky checkout test

Byte · Codex · acme-shop / fix-the-flaky-checkout-test-ca81

Stop

Terminal

Changes

Log

IT ASKS

Update the snapshot or fix the selector?

Type your answer...

Send

Codex (demo stand-in: no model runs here)

> Fix the flaky checkout test

• Read(web/checkout.test.js)

• Bash(npm test -- web)

not ok 1 - checkout shows the total

failed 2 of 5 runs: .total was empty

• Read(web/checkout.js)

.total is drawn twice and the first one is empty while the order loads.

? Update the snapshot or fix the selector?

>

Keys

Demo data: scripted agents, no model runs


Recording: about 60 s, silent, the real build in demo mode with scripted stand-in agents / tests/demo.test.ts, tests/demo-e2e.test.ts

KIPDECK · INVESTOR DECK · DATA AS OF 9 OCT 2026

06 / 15

Self-hosted, vendor-neutral, on your own machine.

Kipdeck itself is self-hosted: no Kipdeck cloud, no account. Your agents still talk to their own model APIs.



- RUNS FROM SOURCE TODAY; NPM AT M1
- NO PASSWORD ON YOUR OWN MACHINE
- NO KIPDECK CLOUD; AGENTS USE THEIR OWN APIS

— PROOF OF EXECUTION

A working product, measured. No users yet.

FRESH START TO FIRST MERGED CHANGE

10.7s



4 CLICKS · MEASURED 2026-10-07, HEADLESS, STAND-IN AGENT, NO MODEL

3.4s

to the first agent at work,
1 click

1,528

tests in the suite, run by
npm test

19

security audit findings
closed before launch

Next, before revenue: publish, launch, and sign 5 design-partner teams that run 5+ agents a day.

Times and clicks: design/measure-final.mjs, run 2026-10-07 (it writes measure.json) / Tests: npm test lists 1,528, 2026-10-09 / Security audit: internal review, 2026-10-06


— MARKET, BOTTOM-UP

EACH DOT = 10,000 DEVELOPERS

11.3M developers use agents.


Professional developers SlashData, early 2025	36.5M
x 31% actively use AI agents Stack Overflow 2025 share; our estimate	11.3M
x USD 360 a seat a year = TAM our team-tier price; at Conductor's USD 720: 8.1B	USD 4.1B
Daily agent users, 14.1% = SAM Stack Overflow 2025; 5.1M x USD 360	USD 1.9B
Year-5 path, our assumption 100k free installs x 2% become teams x 10 seats x USD 360	USD 7.2M ARR

The budget already exists: USD 5.5B+ a year already goes to AI coding tools (slide 4).



COMPETITION

Conductor and Superset help one developer run more agents. We stop your team keeping them waiting.



Distribution first. Then wait data no single vendor has.

01 **Open-source installs**

Single-player is MIT-licensed and free; the repository is public today, npm publish at M1.

02 **Cross-vendor wait data**

Every merge already records agent, model, reviewer and wait. As teams adopt, that becomes data no single vendor can collect.

03 **Team upgrade**

The second person who needs the queue turns a free install into a paid team.

SIGNED MERGE RECORDS

ILLUSTRATION


#0412 ✓ signed
Codex
waited 4m

#0413 ✓ signed
Claude Code
waited 1m

#0414 ✓ signed
Cursor
waited 9m

MEDIAN WAIT ON A PERSON, BY AGENT, WEEK OVER WEEK

MOCK DATA



 Self-hosted: the record sits on the team's machine, not in a Kipdeck cloud. The code is not the moat; installs and the wait data are.

Free for one engineer. Teams pay.

SINGLE PLAYER Free MIT Inbox, local merge, signed log, wait metrics The distribution: public today; one npx command at M1	TEAM TIER PRICE IS OUR ASSUMPTION TEAM USD 30 /seat/mo Shared queue, routing to the owner, team wait-time reports SSO, audit export. Not built yet: milestone M3
--	---

GO-TO-MARKET

01 Launch npm publish and Hacker News with the 60 s demo (M1).	02 5 design partners Founder-led, teams running 5+ agents a day; European and regulated first.	03 Bottom-up One engineer installs free; the team upgrades when a second person needs the queue.	04 Team tier Paid seats for design partners first, then self-serve.
---	---	---	--

No revenue yet. Paid design-partner pilots are possible (offer ready, none sold); that is design-partner revenue, not the business.

— THE TEAM

Engineers who have built and scaled products before.



Žygimantas Bagdzevičius

CEO / CTO · [LINKEDIN](#)

400k MAU, EXITED

Ex-Amazon engineer. Founded and exited a Lithuanian business that scaled to 400k monthly active users. Built Motored's whole product before outside money.



Lukas Kveraga

FOUNDING ENGINEER

Ex-Vinted

Lithuania's marketplace unicorn. Hands-on experience scaling a high-traffic two-sided marketplace.



Ernestas Rimkevičius

SOFTWARE ENGINEER · [LINKEDIN](#)

Pipelines


Backend, importer, listings pipeline and ops tooling (at Motored).

DISCLOSURE AND COMMITMENT

Today Žygimantas is also co-founder and CEO/CTO of Motored (Antler-backed, raising its own pre-seed), and Lukas and Ernestas are engineers there. Who goes full-time on Kipdeck and from when, the equity split and IP assignment to the Kipdeck company are agreed in writing before this round closes and shared in diligence.

Company: to be incorporated before close.

CEO/CTO COMMITS WITH AGENTS, 30 SEP-9 OCT




motored-lt.vercel.app/data-room (backgrounds and Motored roles) / 570 commits by the CEO/CTO since 2026-09-30 (main at 9a30f41f, 2026-10-09), almost all co-authored with Claude; not counted: 2 Dependabot bumps, 2 upstream snapshot imports, 11 re-committed upstream PRs


THE ASK

USD 100k

Post-money SAFE, USD 1M valuation cap: 10%.
Company to be incorporated before close. About USD 11k a month: 9 months of runway.



9 MONTHS TO 5 DESIGN PARTNERS AND THE FIRST PAYING TEAMS



NEEDS YOU 0			
0c	Upgrade the payment SDK to v5 OpenCode	MERGED	waited 34m
cc	Add rate limiting to /api/login Claude Code	MERGED	waited 8m
Mu	Bump the Node image to 22 Muse	MERGED	waited 1m

Nobody waiting.
What this round proves: teams pay for one queue across vendors.

Žygimantas Bagdzevičius, Lukas Kveraga, Ernestas Rimkevičius / github.com/zbagdzevicius/kipdeck

Tool by tool, and more evidence on the problem.

WHAT YOU NEED	VENDOR DASHBOARD	TASK BOARD	GITHUB PRS	TERMINALS, TMUX	KIPDECK
Every agent CLI in one list	Its own	If you file them	Once a PR exists	One pane at a time	Ranked by wait
Knows an agent is blocked on you	Its own	No	No	If you look	Top of the list
Answer in one box	In its app	No	No	Find the terminal	Every CLI
Review and merge before a PR	No	No	No	By hand	Or open one
Measures wait on a person	No	No	No	No	Per agent, model
Self-hosted	Varies (local CLI or vendor cloud)	Depends	Their cloud	Yes	Yes, MIT

METR: experienced devs were 19% slower with AI while believing they were 20% faster.

Stack Overflow 2025: 66% frustrated by "almost right" output; 3.1% highly trust AI.

DORA 2025: AI adoption has a negative relationship with delivery stability.

Atlassian DevEx, 3,500 devs: 50% lose 10+ hours a week to friction; 63% say leaders do not get it.